

# Zhanhui Zhou

asap.zzhou@gmail.com / 86-186-2154-1215 / github/ZHZisZZ

## Education

---

**University of Michigan (dual degree)**, Ann Arbor, MI 2020 - 2022  
B.S.E., Computer Science (Highest Honors)  
GPA: 4.00/4.00

**Shanghai Jiao Tong University (dual degree)**, Shanghai, China 2018 - 2020  
B.S.E., Computer Engineering (Highest Honors)  
GPA: 3.78/4.00

## Professional Experience

---

**Shanghai AI Lab (AI Alignment Team)**, Shanghai, China 2022 - Now  
Research Engineer, full-time  
○ Lead research on AI alignment and LLMs.

**Shanghai AI Lab (Embodied-AI Team)**, Shanghai, China 2022  
Research Intern, part-time  
○ Lead large-scale distributed reinforcement learning infrastructure.

## Selected Publications

---

\* indicates equal contribution

**Zhanhui Zhou**, Zhixuan Liu, Jie Liu, Zhichen Dong, Chao Yang, Yu Qiao. 2024  
Weak-to-Strong Search: Align Large Language Models via Searching over Small Language Models. *Preprint, under review.*

**Zhanhui Zhou**, Jie Liu, Zhichen Dong, Jiaheng Liu, Chao Yang, Wanli Ouyang, Yu Qiao. 2024  
Emulated Disalignment: Safety Alignment for Large Language Models May Backfire! **Outstanding Paper Award (1% of submission)**. *Annual Meeting of the Association for Computational Linguistics (ACL)*.

**Zhanhui Zhou\***, Jie Liu\*, Chao Yang, Jing Shao, Yu Liu, Xiangyu Yue, Wanli Ouyang, Yu Qiao. 2024  
Beyond One-Preference-Fits-All Alignment: Multi-Objective Direct Preference Optimization. *Findings of the Association for Computational Linguistics ACL (Findings of ACL)*.

**Zhanhui Zhou\***, Man To Tang\*, Qiping Pan\*, Shangyin Tan, Xinyu Wang, Tianyi Zhang. 2022  
INTENT: Interactive TENSor Transformation Synthesis. *Symposium on User Interface Software and Technology (UIST)*.

## Other Publications

---

\* indicates equal contribution

Jie Liu*, <b>Zhanhui Zhou*</b> , Jiaheng Liu, Xingyuan Bu, Chao Yang, Han-Sen Zhong, Wanli Ouyang. Iterative Length-Regularized Direct Preference Optimization: A Case Study on Improving 7B Language Models to GPT-4 Level. <i>Preprint, under review.</i>	2024
Ge Bai, Jie Liu, Xingyuan Bu, Yancheng He, Jiaheng Liu, <b>Zhanhui Zhou</b> , Zhuoran Lin, Wenbo Su, Tiezheng Ge, Bo Zheng, Wanli Ouyang. MT-Bench-101: A Fine-Grained Benchmark for Evaluating Large Language Models in Multi-Turn Dialogues. <i>Annual Meeting of the Association for Computational Linguistics (ACL).</i>	2024
Yanan Wu, Jie Liu, Xingyuan Bu, Jiaheng Liu, <b>Zhanhui Zhou</b> , Yuanxing Zhang, Chenchen Zhang, Zhiqi Bai, Haibin Chen, Tiezheng Ge, Wanli Ouyang, Wenbo Su, Bo Zheng. ConceptMath: A Bilingual Concept-wise Benchmark for Measuring Mathematical Reasoning of Large Language Models. <i>Findings of the Association for Computational Linguistics ACL (Findings of ACL).</i>	2024
Zhichen Dong*, <b>Zhanhui Zhou*</b> , Chao Yang, Jing Shao, Yu Qiao. Attacks, Defenses and Evaluations for LLM Conversation Safety: A Survey. <i>Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL).</i>	2024

## Academic Awards

---

College of Engineering - Dean's Honor List (Umich)	All semesters
Second-Class Academic Excellence Scholarship (top 10%) (SJTU)	All semesters

## Service

---

NeurIPS Conference Reviewer	2024
ICLR Conference Reviewer	2024

## Skills

---

Language: English (fluent), Chinese (native)

Programming Language: Python, C++, C

Skills: PyTorch, TensorFlow,  $\text{\LaTeX}$